


# Artificial Neural Network (ANN) VS Convolutional Neural Network (CNN)


Using both **ANN** and **CNN** on **Digits Mnist dataset**

[https://en.wikipedia.org/wiki/MNIST\\_database](https://en.wikipedia.org/wiki/MNIST_database)

Artificial Neural Network (ANN)



Convolutional Neural Network (CNN)



## Artificial Neural Network (ANN)

```
In [1]: # !pip install tensorflow
```

```
In [47]: # Importing required Libraries
import warnings
warnings.filterwarnings("ignore")
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import tensorflow as tf
import os
import seaborn as sns
import keras
from keras.models import Sequential
from keras.layers import Dense, Flatten, BatchNormalization, Dropout
from keras.layers import LeakyReLU, Conv2D, MaxPooling2D, AveragePooling2D
from keras import regularizers
import time
```

```
In [3]: print(tf.__version__) # Tensorflow version
2.13.0
```

```
In [4]: tf.config.list_physical_devices("GPU") # Finding GPU
```

```
Out[4]: []
```

```
In [5]: tf.config.list_physical_devices("CPU") # Finding CPU
```

```
Out[5]: [PhysicalDevice(name='/physical_device:CPU:0', device_type='CPU')]
```


```
In [6]: # Loading Mnist dataset from tensorflow
mnist = tf.keras.datasets.mnist
(x_train_full, y_train_full) , (x_test, y_test) = mnist.load_data()
```

```
In [7]: # Shape of Mnist dataset
print(x_train_full.shape) # Training data shape
print(x_test.shape)      # Testing data shape
```

```
(60000, 28, 28)
(10000, 28, 28)
```

```
In [8]: # There are 60000 image of handwritten digits and each image has resolution 28*28 which is equal to 784 pixels
```

```
In [9]: # Here is some example
k = 5
plt.figure(figsize=(15,5))
for i in range(1,k+1):
    plt.subplot(1,k,i)
    idx = int(np.random.randint(0,60000,1))
    plt.imshow(x_train_full[idx], cmap="binary")
plt.show()
```



```
In [10]: x_train_full[0] # Image at "0" index
```

```

Out[10]: array([[ 0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,
                  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,
                  0,  0],
                [ 0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,
                  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,
                  0,  0],
                [ 0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,
                  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,
                  0,  0],
                [ 0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,
                  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,
                  0,  0],
                [ 0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,
                  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,
                  0,  0],
                [ 0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  3,
                  18, 18, 18, 126, 136, 175, 26, 166, 255, 247, 127,  0,  0,
                  0,  0],
                [ 0,  0,  0,  0,  0,  0,  0,  0,  30, 36, 94, 154, 170,
                  253, 253, 253, 253, 253, 225, 172, 253, 242, 195, 64,  0,  0,
                  0,  0],
                [ 0,  0,  0,  0,  0,  0,  0,  49, 238, 253, 253, 253, 253,
                  253, 253, 253, 253, 251, 93, 82, 82, 56, 39,  0,  0,  0,
                  0,  0],
                [ 0,  0,  0,  0,  0,  0,  0,  18, 219, 253, 253, 253, 253,
                  253, 198, 182, 247, 241,  0,  0,  0,  0,  0,  0,  0,  0,
                  0,  0],
                [ 0,  0,  0,  0,  0,  0,  0,  0,  80, 156, 107, 253, 253,
                  205, 11,  0,  43, 154,  0,  0,  0,  0,  0,  0,  0,  0,
                  0,  0],
                [ 0,  0,  0,  0,  0,  0,  0,  0,  0,  14,  1, 154, 253,
                  90,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,
                  0,  0],
                [ 0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0, 139, 253,
                  190,  2,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,
                  0,  0],
                [ 0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0, 11, 190,
                  253, 70,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,
                  0,  0],
                [ 0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0, 35,
                  241, 225, 160, 108,  1,  0,  0,  0,  0,  0,  0,  0,  0,
                  0,  0],
                [ 0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,
                  81, 240, 253, 253, 119, 25,  0,  0,  0,  0,  0,  0,  0,
                  0,  0],
                [ 0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,
                  0,  45, 186, 253, 253, 150, 27,  0,  0,  0,  0,  0,  0,
                  0,  0],
                [ 0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,
                  0,  0,  16, 93, 252, 253, 187,  0,  0,  0,  0,  0,  0,
                  0,  0],
                [ 0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,
                  0,  0,  0,  0,  249, 253, 249, 64,  0,  0,  0,  0,  0,
                  0,  0],
                [ 0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,
                  0,  46, 130, 183, 253, 253, 207,  2,  0,  0,  0,  0,  0,
                  0,  0],
                [ 0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0, 39,
                  148, 229, 253, 253, 253, 250, 182,  0,  0,  0,  0,  0,  0,
                  0,  0],
                [ 0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  24, 114, 221,
                  253, 253, 253, 253, 201, 78,  0,  0,  0,  0,  0,  0,
                  0,  0],
                [ 0,  0,  0,  0,  0,  0,  0,  23, 66, 213, 253, 253,
                  253, 253, 198, 81,  2,  0,  0,  0,  0,  0,  0,  0,  0,
                  0,  0],
                [ 0,  0,  0,  0,  0,  18, 171, 219, 253, 253, 253, 253,
                  195, 80,  9,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,
                  0,  0],
                [ 0,  0,  0,  0,  55, 172, 226, 253, 253, 253, 253, 244, 133,
                  11,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,
                  0,  0],
                [ 0,  0,  0,  0, 136, 253, 253, 253, 212, 135, 132, 16,  0,
                  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,
                  0,  0],
                [ 0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,
                  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,
                  0,  0],
                [ 0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,
                  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,
                  0,  0],
                [ 0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,
                  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,
                  0,  0]], dtype=uint8)

```

```
In [11]: y_train_full[0] # Number at "0" index
```

```
In [12]: # Converting the pixels value between 0 to 1
x_valid, x_train = x_train_full[:5000]/255, x_train_full[5000:]/255
y_valid, y_train = y_train_full[:5000], y_train_full[5000:]

x_test = x_test/255
```

```
# Defining the layers
LAYERS = [
    tf.keras.layers.Flatten(input_shape=[28,28],name="inputLayer"),
    tf.keras.layers.Dense(300,activation="relu",name="hiddenLayer1"),
    tf.keras.layers.Dense(100,activation="relu",name="hiddenLayer2"),
    tf.keras.layers.Dense(10,activation="softmax",name="outputLayer")
]

model = tf.keras.models.Sequential(LAYERS)
```

```

Model: "sequential"

Layer (type)                Output Shape                Param #
=====
inputLayer (Flatten)        (None, 784)                0
hiddenLayer1 (Dense)        (None, 300)                235500
hiddenLayer2 (Dense)        (None, 100)                30100
outputLayer (Dense)         (None, 10)                 1010
=====
Total params: 266610 (1.02 MB)
Trainable params: 266610 (1.02 MB)
Non-trainable params: 0 (0.00 Byte)

```

```
Out[15]: []
```

```
In [17]: hidden1 = model.layers[1]
hidden1.get_weights() # The weights of first hidden layer
```

```
In [18]: # Defining Loss_function, Optimizer, Metrics and compiling the model
loss_function = "sparse_categorical_crossentropy"
```

```
OPTIMIZER = tf.keras.optimizers.SGD(learning_rate=0.001)
METRICS = ["accuracy"]

model.compile(
    loss=loss_function,
    optimizer=OPTIMIZER,
    metrics=METRICS
)
```

```
In [19]: # Creating a function for saving the logs at proper folder
def get_log_path(log_dir="logs/fit"):
    filename = time.strftime("1_log_%y_%m_%d_%H_%M_%S")
    logs_path = os.path.join(log_dir, filename)
    print(f"Saving logs at {logs_path}")
    return logs_path
```

```
In [20]: # Creating logs callback
log_dirs = get_log_path()
tb_cb = tf.keras.callbacks.TensorBoard(log_dir=log_dirs)
```

Saving logs at logs/fit/1\_log\_23\_09\_26\_19\_48\_02

```
In [21]: # Creating early stopping callback
early_stopping_cb = tf.keras.callbacks.EarlyStopping(patience=7, restore_best_weights=True)
```

```
In [22]: # Saving the model with callback
CKPT_path = os.path.join("Models", "Model_ckpt_Digit_mnist_1.h5")
checkpoint_cb = tf.keras.callbacks.ModelCheckpoint(CKPT_path, save_best_only=True)
```

```
In [23]: EPOCHS = 50 # Number of Epochs
VALIDATION_SET = (x_valid, y_valid) # Validation data
history = model.fit(x_train, y_train, epochs=EPOCHS, validation_data=VALIDATION_SET, batch_size=64, # Training the model
                    callbacks=[tb_cb, early_stopping_cb, checkpoint_cb], use_multiprocessing=True)
```

Epoch 1/50  
860/860 [=====] - 4s 4ms/step - loss: 2.0446 - accuracy: 0.4097 - val\_loss: 1.7333 - val\_accuracy: 0.6436  
Epoch 2/50  
860/860 [=====] - 3s 4ms/step - loss: 1.4550 - accuracy: 0.7157 - val\_loss: 1.1851 - val\_accuracy: 0.7628  
Epoch 3/50  
860/860 [=====] - 4s 4ms/step - loss: 1.0196 - accuracy: 0.7898 - val\_loss: 0.8604 - val\_accuracy: 0.8108  
Epoch 4/50  
860/860 [=====] - 3s 4ms/step - loss: 0.7823 - accuracy: 0.8234 - val\_loss: 0.6892 - val\_accuracy: 0.8376  
Epoch 5/50  
860/860 [=====] - 3s 4ms/step - loss: 0.6523 - accuracy: 0.8436 - val\_loss: 0.5889 - val\_accuracy: 0.8564  
Epoch 6/50  
860/860 [=====] - 3s 4ms/step - loss: 0.5727 - accuracy: 0.8573 - val\_loss: 0.5238 - val\_accuracy: 0.8694  
Epoch 7/50  
860/860 [=====] - 3s 4ms/step - loss: 0.5191 - accuracy: 0.8671 - val\_loss: 0.4781 - val\_accuracy: 0.8798  
Epoch 8/50  
860/860 [=====] - 3s 4ms/step - loss: 0.4806 - accuracy: 0.8746 - val\_loss: 0.4446 - val\_accuracy: 0.8860  
Epoch 9/50  
860/860 [=====] - 4s 4ms/step - loss: 0.4515 - accuracy: 0.8803 - val\_loss: 0.4181 - val\_accuracy: 0.8916  
Epoch 10/50  
860/860 [=====] - 3s 4ms/step - loss: 0.4287 - accuracy: 0.8844 - val\_loss: 0.3975 - val\_accuracy: 0.8954  
Epoch 11/50  
860/860 [=====] - 3s 4ms/step - loss: 0.4102 - accuracy: 0.8883 - val\_loss: 0.3804 - val\_accuracy: 0.8988  
Epoch 12/50  
860/860 [=====] - 3s 4ms/step - loss: 0.3947 - accuracy: 0.8915 - val\_loss: 0.3666 - val\_accuracy: 0.9014  
Epoch 13/50  
860/860 [=====] - 3s 4ms/step - loss: 0.3817 - accuracy: 0.8940 - val\_loss: 0.3545 - val\_accuracy: 0.9036  
Epoch 14/50  
860/860 [=====] - 3s 4ms/step - loss: 0.3705 - accuracy: 0.8967 - val\_loss: 0.3441 - val\_accuracy: 0.9052  
Epoch 15/50  
860/860 [=====] - 3s 4ms/step - loss: 0.3607 - accuracy: 0.8986 - val\_loss: 0.3350 - val\_accuracy: 0.9072  
Epoch 16/50  
860/860 [=====] - 3s 4ms/step - loss: 0.3520 - accuracy: 0.9012 - val\_loss: 0.3265 - val\_accuracy: 0.9082  
Epoch 17/50  
860/860 [=====] - 3s 4ms/step - loss: 0.3441 - accuracy: 0.9027 - val\_loss: 0.3193 - val\_accuracy: 0.9094  
Epoch 18/50  
860/860 [=====] - 3s 4ms/step - loss: 0.3369 - accuracy: 0.9049 - val\_loss: 0.3126 - val\_accuracy: 0.9116  
Epoch 19/50  
860/860 [=====] - 3s 4ms/step - loss: 0.3304 - accuracy: 0.9063 - val\_loss: 0.3067 - val\_accuracy: 0.9148  
Epoch 20/50  
860/860 [=====] - 3s 4ms/step - loss: 0.3243 - accuracy: 0.9083 - val\_loss: 0.3010 - val\_accuracy: 0.9158  
Epoch 21/50  
860/860 [=====] - 3s 4ms/step - loss: 0.3187 - accuracy: 0.9096 - val\_loss: 0.2957 - val\_accuracy: 0.9168  
Epoch 22/50  
860/860 [=====] - 3s 4ms/step - loss: 0.3135 - accuracy: 0.9110 - val\_loss: 0.2911 - val\_accuracy: 0.9182  
Epoch 23/50  
860/860 [=====] - 3s 4ms/step - loss: 0.3087 - accuracy: 0.9122 - val\_loss: 0.2862 - val\_accuracy: 0.9202  
Epoch 24/50  
860/860 [=====] - 3s 4ms/step - loss: 0.3039 - accuracy: 0.9135 - val\_loss: 0.2823 - val\_accuracy: 0.9210  
Epoch 25/50  
860/860 [=====] - 3s 4ms/step - loss: 0.2995 - accuracy: 0.9143 - val\_loss: 0.2782 - val\_accuracy: 0.9216  
Epoch 26/50  
860/860 [=====] - 3s 4ms/step - loss: 0.2954 - accuracy: 0.9157 - val\_loss: 0.2742 - val\_accuracy: 0.9234  
Epoch 27/50  
860/860 [=====] - 3s 4ms/step - loss: 0.2913 - accuracy: 0.9168 - val\_loss: 0.2707 - val\_accuracy: 0.9236  
Epoch 28/50  
860/860 [=====] - 3s 4ms/step - loss: 0.2875 - accuracy: 0.9180 - val\_loss: 0.2670 - val\_accuracy: 0.9246  
Epoch 29/50  
860/860 [=====] - 3s 4ms/step - loss: 0.2839 - accuracy: 0.9192 - val\_loss: 0.2637 - val\_accuracy: 0.9256

```

Epoch 30/50
860/860 [=====] - 3s 4ms/step - loss: 0.2804 - accuracy: 0.9205 - val_loss: 0.2607 - val_accuracy: 0.9256
Epoch 31/50
860/860 [=====] - 3s 4ms/step - loss: 0.2769 - accuracy: 0.9206 - val_loss: 0.2572 - val_accuracy: 0.9272
Epoch 32/50
860/860 [=====] - 3s 4ms/step - loss: 0.2737 - accuracy: 0.9216 - val_loss: 0.2545 - val_accuracy: 0.9270
Epoch 33/50
860/860 [=====] - 3s 4ms/step - loss: 0.2706 - accuracy: 0.9228 - val_loss: 0.2519 - val_accuracy: 0.9280
Epoch 34/50
860/860 [=====] - 3s 4ms/step - loss: 0.2675 - accuracy: 0.9241 - val_loss: 0.2489 - val_accuracy: 0.9286
Epoch 35/50
860/860 [=====] - 3s 4ms/step - loss: 0.2645 - accuracy: 0.9249 - val_loss: 0.2462 - val_accuracy: 0.9304
Epoch 36/50
860/860 [=====] - 3s 4ms/step - loss: 0.2617 - accuracy: 0.9256 - val_loss: 0.2433 - val_accuracy: 0.9312
Epoch 37/50
860/860 [=====] - 3s 4ms/step - loss: 0.2588 - accuracy: 0.9266 - val_loss: 0.2409 - val_accuracy: 0.9320
Epoch 38/50
860/860 [=====] - 3s 4ms/step - loss: 0.2561 - accuracy: 0.9271 - val_loss: 0.2385 - val_accuracy: 0.9328
Epoch 39/50
860/860 [=====] - 3s 4ms/step - loss: 0.2534 - accuracy: 0.9278 - val_loss: 0.2356 - val_accuracy: 0.9338
Epoch 40/50
860/860 [=====] - 3s 4ms/step - loss: 0.2509 - accuracy: 0.9284 - val_loss: 0.2337 - val_accuracy: 0.9336
Epoch 41/50
860/860 [=====] - 3s 4ms/step - loss: 0.2484 - accuracy: 0.9298 - val_loss: 0.2312 - val_accuracy: 0.9348
Epoch 42/50
860/860 [=====] - 3s 4ms/step - loss: 0.2459 - accuracy: 0.9301 - val_loss: 0.2291 - val_accuracy: 0.9364
Epoch 43/50
860/860 [=====] - 3s 4ms/step - loss: 0.2435 - accuracy: 0.9306 - val_loss: 0.2270 - val_accuracy: 0.9362
Epoch 44/50
860/860 [=====] - 3s 4ms/step - loss: 0.2412 - accuracy: 0.9315 - val_loss: 0.2247 - val_accuracy: 0.9372
Epoch 45/50
860/860 [=====] - 3s 4ms/step - loss: 0.2389 - accuracy: 0.9321 - val_loss: 0.2227 - val_accuracy: 0.9374
Epoch 46/50
860/860 [=====] - 3s 4ms/step - loss: 0.2367 - accuracy: 0.9326 - val_loss: 0.2211 - val_accuracy: 0.9382
Epoch 47/50
860/860 [=====] - 3s 4ms/step - loss: 0.2345 - accuracy: 0.9335 - val_loss: 0.2188 - val_accuracy: 0.9384
Epoch 48/50
860/860 [=====] - 3s 4ms/step - loss: 0.2323 - accuracy: 0.9338 - val_loss: 0.2169 - val_accuracy: 0.9384
Epoch 49/50
860/860 [=====] - 3s 4ms/step - loss: 0.2302 - accuracy: 0.9343 - val_loss: 0.2151 - val_accuracy: 0.9400
Epoch 50/50
860/860 [=====] - 3s 4ms/step - loss: 0.2282 - accuracy: 0.9348 - val_loss: 0.2134 - val_accuracy: 0.9400

```

```
In [24]: pd.DataFrame(history.history) # All the losses and accuracy in each epoch
```

Out[24]:


	loss	accuracy	val_loss	val_accuracy
0	2.044599	0.409745	1.733334	0.6436
1	1.455023	0.715709	1.185071	0.7628
2	1.019633	0.789800	0.860445	0.8108
3	0.782261	0.823436	0.689154	0.8376
4	0.652308	0.843582	0.588874	0.8564
5	0.572710	0.857255	0.523771	0.8694
6	0.519132	0.867127	0.478128	0.8798
7	0.480602	0.874636	0.444642	0.8860
8	0.451508	0.880273	0.418126	0.8916
9	0.428673	0.884400	0.397529	0.8954
10	0.410156	0.888255	0.380364	0.8988
11	0.394743	0.891545	0.366589	0.9014
12	0.381738	0.894036	0.354540	0.9036
13	0.370523	0.896655	0.344144	0.9052
14	0.360709	0.898564	0.334979	0.9072
15	0.351955	0.901182	0.326504	0.9082
16	0.344110	0.902709	0.319325	0.9094
17	0.336940	0.904909	0.312572	0.9116
18	0.330399	0.906345	0.306733	0.9148
19	0.324309	0.908327	0.300988	0.9158
20	0.318691	0.909636	0.295740	0.9168
21	0.313513	0.911000	0.291094	0.9182
22	0.308652	0.912200	0.286233	0.9202
23	0.303929	0.913527	0.282320	0.9210
24	0.299470	0.914255	0.278207	0.9216
25	0.295367	0.915745	0.274233	0.9234
26	0.291341	0.916800	0.270697	0.9236
27	0.287510	0.917964	0.267022	0.9246
28	0.283905	0.919200	0.263742	0.9256
29	0.280367	0.920509	0.260654	0.9256
30	0.276936	0.920564	0.257197	0.9272
31	0.273685	0.921582	0.254464	0.9270
32	0.270556	0.922836	0.251917	0.9280
33	0.267517	0.924127	0.248867	0.9286
34	0.264475	0.924945	0.246232	0.9304
35	0.261701	0.925564	0.243337	0.9312
36	0.258824	0.926582	0.240950	0.9320
37	0.256124	0.927091	0.238550	0.9328
38	0.253415	0.927782	0.235629	0.9338
39	0.250865	0.928364	0.233744	0.9336
40	0.248369	0.929782	0.231237	0.9348
41	0.245898	0.930109	0.229094	0.9364
42	0.243542	0.930636	0.227009	0.9362
43	0.241214	0.931473	0.224715	0.9372
44	0.238904	0.932091	0.222734	0.9374
45	0.236651	0.932600	0.221052	0.9382
46	0.234489	0.933545	0.218850	0.9384
47	0.232283	0.933782	0.216922	0.9384
48	0.230208	0.934345	0.215123	0.9400
49	0.228181	0.934836	0.213441	0.9400

In [25]:

pd.DataFrame(history.history).plot() # Plotting all of the accuracy and losses



Out[25]: <AxesSubplot: >



```
In [26]: ckpt_model = tf.keras.models.load_model(CKPT_path) # Loading model
         ckpt_model.evaluate(x_test,y_test)                # Evaluating the performance of the model on test data
```

313/313 [=====] - 1s 2ms/step - loss: 0.2221 - accuracy: 0.9365

Out[26]: [0.2220998853445053, 0.9365000128746033]

The test accuracy of the first **ANN** model is **93.65 %**

```
In [27]: # Predicting the probability of 5 samples randomly
print("Probability of each number for image of testing data:- \n")
for j in range(5):
    idx = np.random.randint(0,10000,1)[0]
    y_prob_lst = ckpt_model.predict(x_test[idx-1:idx]).round(3)[0]
    print("\nNumber : Probability")
    for i in range(10):
        print(i, " : ", y_prob_lst[i])
    y_predict = np.argmax(y_prob_lst,axis=-1)
    print("Predicted final output is: ",y_predict,"\n")
```

Probability of each number for image of testing data:-

1/1 [=====] - 0s 105ms/step

Number : Probability

0	:	0.0
1	:	0.0
2	:	0.001
3	:	0.001
4	:	0.05
5	:	0.0
6	:	0.0
7	:	0.007
8	:	0.001
9	:	0.94

Predicted final output is: 9

1/1 [=====] - 0s 26ms/step

Number : Probability

0	:	0.289
1	:	0.0
2	:	0.005
3	:	0.0
4	:	0.439
5	:	0.048
6	:	0.035
7	:	0.008
8	:	0.121
9	:	0.054

Predicted final output is: 4

1/1 [=====] - 0s 25ms/step

Number : Probability

0	:	0.002
1	:	0.0
2	:	0.002
3	:	0.006
4	:	0.002
5	:	0.697
6	:	0.0
7	:	0.019
8	:	0.026
9	:	0.245

Predicted final output is: 5

1/1 [=====] - 0s 26ms/step

Number : Probability

0	:	0.0
1	:	0.0
2	:	0.0
3	:	0.003
4	:	0.051
5	:	0.001
6	:	0.0
7	:	0.047
8	:	0.001
9	:	0.896

Predicted final output is: 9

1/1 [=====] - 0s 26ms/step

Number : Probability

0	:	0.0
1	:	0.0
2	:	0.999
3	:	0.0
4	:	0.0
5	:	0.0
6	:	0.0
7	:	0.0
8	:	0.001
9	:	0.0

Predicted final output is: 2

```
In [28]: del model
del ckpt_model
```

## Second ANN Network

```
In [29]: # Added kernel regularization, batch normalization, dropout and leakyrelu activation function
model = Sequential()
model.add(Flatten(input_shape=[28,28]))
```

```

model.add(Dense(units=256,kernel_regularizer=regularizers.L1L2(0.0001,0.0001)))
model.add(BatchNormalization())
model.add(LeakyReLU())
model.add(Dropout(0.1))
model.add(Dense(units=128,kernel_regularizer=regularizers.L1L2(0.0001,0.0001)))
model.add(BatchNormalization())
model.add(LeakyReLU())
model.add(Dropout(0.1))
model.add(Dense(units=64,kernel_regularizer=regularizers.L1L2(0.0001,0.0001)))
model.add(BatchNormalization())
model.add(LeakyReLU())
model.add(Dropout(0.1))
model.add(Dense(units=32,kernel_regularizer=regularizers.L1L2(0.0001,0.0001)))
model.add(BatchNormalization())
model.add(LeakyReLU())
model.add(Dropout(0.1))
model.add(Dense(units=16,kernel_regularizer=regularizers.L1L2(0.0001,0.0001)))
model.add(BatchNormalization())
model.add(LeakyReLU())
model.add(Dropout(0.1))
model.add(Dense(units=10,activation="softmax"))

```

In [30]: `model.summary()`

Model: "sequential\_1"

Layer (type)	Output Shape	Param #
=====		
flatten (Flatten)	(None, 784)	0
dense (Dense)	(None, 256)	200960
batch_normalization (Batch Normalization)	(None, 256)	1024
leaky_re_lu (LeakyReLU)	(None, 256)	0
dropout (Dropout)	(None, 256)	0
dense_1 (Dense)	(None, 128)	32896
batch_normalization_1 (Batch Normalization)	(None, 128)	512
leaky_re_lu_1 (LeakyReLU)	(None, 128)	0
dropout_1 (Dropout)	(None, 128)	0
dense_2 (Dense)	(None, 64)	8256
batch_normalization_2 (Batch Normalization)	(None, 64)	256
leaky_re_lu_2 (LeakyReLU)	(None, 64)	0
dropout_2 (Dropout)	(None, 64)	0
dense_3 (Dense)	(None, 32)	2080
batch_normalization_3 (Batch Normalization)	(None, 32)	128
leaky_re_lu_3 (LeakyReLU)	(None, 32)	0
dropout_3 (Dropout)	(None, 32)	0
dense_4 (Dense)	(None, 16)	528
batch_normalization_4 (Batch Normalization)	(None, 16)	64
leaky_re_lu_4 (LeakyReLU)	(None, 16)	0
dropout_4 (Dropout)	(None, 16)	0
dense_5 (Dense)	(None, 10)	170
=====		
Total params: 246874 (964.35 KB)		
Trainable params: 245882 (960.48 KB)		
Non-trainable params: 992 (3.88 KB)		

In [31]: `def get_log_path(log_dir="logs/fit"):`  
`filename = time.strftime("2_log_%y_%m_%d_%H_%M_%S")`  
`logs_path = os.path.join(log_dir, filename)`  
`print(f"Saving logs at {logs_path}")`

```
    return logs_path

log_dirs = get_log_path()
tb_cb = tf.keras.callbacks.TensorBoard(log_dir=log_dirs)

early_stopping_cb = tf.keras.callbacks.EarlyStopping(patience=7, restore_best_weights=True)

CKPT_path = os.path.join("Models", "Model_ckpt_Digit_mnist_2.h5")
checkpoint_cb = tf.keras.callbacks.ModelCheckpoint(CKPT_path, save_best_only=True)

EPOCHS = 50
VALIDATION_SET = (x_valid, y_valid)

loss_function = "sparse_categorical_crossentropy"
OPTIMIZER = tf.keras.optimizers.Adam(learning_rate=0.001)
METRICS = ["accuracy"]

model.compile(
    loss=loss_function,
    optimizer=OPTIMIZER,
    metrics=METRICS
)

history = model.fit(x_train, y_train, epochs=EPOCHS, validation_data=VALIDATION_SET, batch_size=64,
                    callbacks=[tb_cb, early_stopping_cb, checkpoint_cb], use_multiprocessing=True)
```

Saving logs at logs/fit/2\_log\_23\_09\_26\_19\_50\_51

Epoch 1/50

860/860 [=====] - 11s 9ms/step - loss: 1.1647 - accuracy: 0.8697 - val\_loss: 0.6543 - val\_accuracy: 0.9488

Epoch 2/50

860/860 [=====] - 7s 8ms/step - loss: 0.6878 - accuracy: 0.9233 - val\_loss: 0.5432 - val\_accuracy: 0.9448

Epoch 3/50

860/860 [=====] - 7s 8ms/step - loss: 0.5877 - accuracy: 0.9297 - val\_loss: 0.4598 - val\_accuracy: 0.9600

Epoch 4/50

860/860 [=====] - 7s 8ms/step - loss: 0.5365 - accuracy: 0.9348 - val\_loss: 0.4354 - val\_accuracy: 0.9560

Epoch 5/50

860/860 [=====] - 7s 8ms/step - loss: 0.5017 - accuracy: 0.9392 - val\_loss: 0.4121 - val\_accuracy: 0.9634

Epoch 6/50

860/860 [=====] - 7s 8ms/step - loss: 0.4760 - accuracy: 0.9425 - val\_loss: 0.4001 - val\_accuracy: 0.9640

Epoch 7/50

860/860 [=====] - 7s 8ms/step - loss: 0.4614 - accuracy: 0.9445 - val\_loss: 0.3945 - val\_accuracy: 0.9588

Epoch 8/50

860/860 [=====] - 7s 8ms/step - loss: 0.4515 - accuracy: 0.9437 - val\_loss: 0.3675 - val\_accuracy: 0.9680

Epoch 9/50

860/860 [=====] - 7s 8ms/step - loss: 0.4310 - accuracy: 0.9468 - val\_loss: 0.3562 - val\_accuracy: 0.9654

Epoch 10/50

860/860 [=====] - 7s 8ms/step - loss: 0.4226 - accuracy: 0.9480 - val\_loss: 0.3484 - val\_accuracy: 0.9682

Epoch 11/50

860/860 [=====] - 7s 8ms/step - loss: 0.4131 - accuracy: 0.9479 - val\_loss: 0.3498 - val\_accuracy: 0.9636

Epoch 12/50

860/860 [=====] - 7s 8ms/step - loss: 0.4058 - accuracy: 0.9489 - val\_loss: 0.3338 - val\_accuracy: 0.9714

Epoch 13/50

860/860 [=====] - 7s 8ms/step - loss: 0.3914 - accuracy: 0.9501 - val\_loss: 0.3233 - val\_accuracy: 0.9696

Epoch 14/50

860/860 [=====] - 7s 8ms/step - loss: 0.3885 - accuracy: 0.9516 - val\_loss: 0.3288 - val\_accuracy: 0.9678

Epoch 15/50

860/860 [=====] - 7s 8ms/step - loss: 0.3823 - accuracy: 0.9516 - val\_loss: 0.3065 - val\_accuracy: 0.9746

Epoch 16/50

860/860 [=====] - 7s 8ms/step - loss: 0.3820 - accuracy: 0.9507 - val\_loss: 0.3134 - val\_accuracy: 0.9716

Epoch 17/50

860/860 [=====] - 7s 8ms/step - loss: 0.3751 - accuracy: 0.9516 - val\_loss: 0.3318 - val\_accuracy: 0.9662

Epoch 18/50

860/860 [=====] - 7s 8ms/step - loss: 0.3694 - accuracy: 0.9523 - val\_loss: 0.2952 - val\_accuracy: 0.9742

Epoch 19/50

860/860 [=====] - 7s 8ms/step - loss: 0.3640 - accuracy: 0.9528 - val\_loss: 0.3038 - val\_accuracy: 0.9704

Epoch 20/50

860/860 [=====] - 7s 8ms/step - loss: 0.3642 - accuracy: 0.9527 - val\_loss: 0.3097 - val\_accuracy: 0.9684

Epoch 21/50

860/860 [=====] - 7s 8ms/step - loss: 0.3625 - accuracy: 0.9533 - val\_loss: 0.2941 - val\_accuracy: 0.9716

Epoch 22/50

860/860 [=====] - 7s 8ms/step - loss: 0.3547 - accuracy: 0.9535 - val\_loss: 0.2917 - val\_accuracy: 0.9734

Epoch 23/50

860/860 [=====] - 7s 8ms/step - loss: 0.3564 - accuracy: 0.9531 - val\_loss: 0.2963 - val\_accuracy: 0.9730

Epoch 24/50

860/860 [=====] - 7s 8ms/step - loss: 0.3505 - accuracy: 0.9555 - val\_loss: 0.2995 - val\_accuracy: 0.9716

Epoch 25/50

860/860 [=====] - 7s 8ms/step - loss: 0.3490 - accuracy: 0.9545 - val\_loss: 0.2903 - val\_accuracy: 0.9706

Epoch 26/50

860/860 [=====] - 7s 8ms/step - loss: 0.3475 - accuracy: 0.9547 - val\_loss: 0.2815 - val\_accuracy: 0.9748

Epoch 27/50

860/860 [=====] - 7s 8ms/step - loss: 0.3446 - accuracy: 0.9557 - val\_loss: 0.2833 - val\_accuracy: 0.9726

Epoch 28/50

860/860 [=====] - 7s 8ms/step - loss: 0.3444 - accuracy: 0.9548 - val\_loss: 0.2782 - val\_accuracy: 0.9750

Epoch 29/50

860/860 [=====] - 7s 8ms/step - loss: 0.3469 - accuracy: 0.9554 - val\_loss: 0.2768 - val\_accuracy:


```

0.9742
Epoch 30/50
860/860 [=====] - 7s 8ms/step - loss: 0.3439 - accuracy: 0.9548 - val_loss: 0.2746 - val_accuracy:
0.9746
Epoch 31/50
860/860 [=====] - 7s 8ms/step - loss: 0.3379 - accuracy: 0.9554 - val_loss: 0.2837 - val_accuracy:
0.9742
Epoch 32/50
860/860 [=====] - 7s 8ms/step - loss: 0.3380 - accuracy: 0.9570 - val_loss: 0.2963 - val_accuracy:
0.9684
Epoch 33/50
860/860 [=====] - 7s 8ms/step - loss: 0.3415 - accuracy: 0.9551 - val_loss: 0.3166 - val_accuracy:
0.9628
Epoch 34/50
860/860 [=====] - 7s 8ms/step - loss: 0.3366 - accuracy: 0.9558 - val_loss: 0.3045 - val_accuracy:
0.9638
Epoch 35/50
860/860 [=====] - 7s 8ms/step - loss: 0.3380 - accuracy: 0.9558 - val_loss: 0.2847 - val_accuracy:
0.9712
Epoch 36/50
860/860 [=====] - 7s 8ms/step - loss: 0.3353 - accuracy: 0.9559 - val_loss: 0.2898 - val_accuracy:
0.9736
Epoch 37/50
860/860 [=====] - 7s 8ms/step - loss: 0.3368 - accuracy: 0.9554 - val_loss: 0.2770 - val_accuracy:
0.9740

```

```
In [32]: pd.DataFrame(history.history).plot()
```

```
Out[32]: <AxesSubplot: >
```



```
In [33]: ckpt_model = tf.keras.models.load_model(CKPT_path)
         ckpt_model.evaluate(x_test,y_test)
```

```
313/313 [=====] - 1s 2ms/step - loss: 0.2719 - accuracy: 0.9740
```

```
Out[33]: [0.27185875177383423, 0.9739999771118164]
```

The accuracy of the second **ANN** model is **97.40 %**

```
In [34]: del model
         del ckpt_model
```

## Third ANN Network

```
In [35]: # Added kernel regularization, batch normalization, dropout, kernel initializer and leakyrelu activation function
model = Sequential()
model.add(Flatten(input_shape=[28,28]))
model.add(Dense(units=256,kernel_regularizer=regularizers.L1L2(0.0001,0.0001),kernel_initializer=tf.keras.initializers.HeNormal
model.add(BatchNormalization())
model.add(LeakyReLU())
model.add(Dropout(0.1))
model.add(Dense(units=128,kernel_regularizer=regularizers.L1L2(0.0001,0.0001),kernel_initializer=tf.keras.initializers.HeNormal
model.add(BatchNormalization())
model.add(LeakyReLU())
model.add(Dropout(0.1))
model.add(Dense(units=64,kernel_regularizer=regularizers.L1L2(0.0001,0.0001),kernel_initializer=tf.keras.initializers.HeNormal
model.add(BatchNormalization())
model.add(LeakyReLU())
model.add(Dropout(0.1))
model.add(Dense(units=32,kernel_regularizer=regularizers.L1L2(0.0001,0.0001),kernel_initializer=tf.keras.initializers.HeNormal

```

```

model.add(BatchNormalization())
model.add(LeakyReLU())
model.add(Dropout(0.1))
model.add(Dense(units=16,kernel_regularizer=regularizers.L1L2(0.0001,0.0001),kernel_initializer=tf.keras.initializers.HeNormal
model.add(BatchNormalization())
model.add(LeakyReLU())
model.add(Dropout(0.1))
model.add(Dense(units=10,activation="softmax"))

```

In [36]: `model.summary()`

Model: "sequential\_2"

Layer (type)	Output Shape	Param #
=====		
flatten_1 (Flatten)	(None, 784)	0
dense_6 (Dense)	(None, 256)	200960
batch_normalization_5 (Batch Normalization)	(None, 256)	1024
leaky_re_lu_5 (LeakyReLU)	(None, 256)	0
dropout_5 (Dropout)	(None, 256)	0
dense_7 (Dense)	(None, 128)	32896
batch_normalization_6 (Batch Normalization)	(None, 128)	512
leaky_re_lu_6 (LeakyReLU)	(None, 128)	0
dropout_6 (Dropout)	(None, 128)	0
dense_8 (Dense)	(None, 64)	8256
batch_normalization_7 (Batch Normalization)	(None, 64)	256
leaky_re_lu_7 (LeakyReLU)	(None, 64)	0
dropout_7 (Dropout)	(None, 64)	0
dense_9 (Dense)	(None, 32)	2080
batch_normalization_8 (Batch Normalization)	(None, 32)	128
leaky_re_lu_8 (LeakyReLU)	(None, 32)	0
dropout_8 (Dropout)	(None, 32)	0
dense_10 (Dense)	(None, 16)	528
batch_normalization_9 (Batch Normalization)	(None, 16)	64
leaky_re_lu_9 (LeakyReLU)	(None, 16)	0
dropout_9 (Dropout)	(None, 16)	0
dense_11 (Dense)	(None, 10)	170
=====		
Total params: 246874 (964.35 KB)		
Trainable params: 245882 (960.48 KB)		
Non-trainable params: 992 (3.88 KB)		

```

In [37]: def get_log_path(log_dir="logs/fit"):
          filename = time.strftime("3_log_%y_%m_%d_%H_%M_%S")
          logs_path = os.path.join(log_dir, filename)
          print(f"Saving logs at {logs_path}")
          return logs_path

log_dirs = get_log_path()
tb_cb = tf.keras.callbacks.TensorBoard(log_dir=log_dirs)

early_stopping_cb = tf.keras.callbacks.EarlyStopping(patience=7, restore_best_weights=True)

CKPT_path = os.path.join("Models", "Model_ckpt_Digit_mnist_3.h5")
checkpoint_cb = tf.keras.callbacks.ModelCheckpoint(CKPT_path, save_best_only=True)

EPOCHS = 50
VALIDATION_SET = (x_valid, y_valid)

```

```
loss_function = "sparse_categorical_crossentropy"
OPTIMIZER = tf.keras.optimizers.Adam(learning_rate=0.001,weight_decay=1e-4) # Weight decay is used here
METRICS = ["accuracy"]

model.compile(
    loss=loss_function,
    optimizer=OPTIMIZER,
    metrics=METRICS
)

history = model.fit(x_train, y_train, epochs=EPOCHS, validation_data=VALIDATION_SET, batch_size=64,
                    callbacks=[tb_cb, early_stopping_cb, checkpoint_cb], use_multiprocessing=True)
```



Saving logs at logs/fit/3\_log\_23\_09\_26\_19\_55\_20

Epoch 1/50

860/860 [=====] - 11s 9ms/step - loss: 1.3140 - accuracy: 0.8640 - val\_loss: 0.7566 - val\_accuracy: 0.9488

Epoch 2/50

860/860 [=====] - 7s 8ms/step - loss: 0.7761 - accuracy: 0.9210 - val\_loss: 0.5952 - val\_accuracy: 0.9506

Epoch 3/50

860/860 [=====] - 7s 8ms/step - loss: 0.6420 - accuracy: 0.9304 - val\_loss: 0.4907 - val\_accuracy: 0.9622

Epoch 4/50

860/860 [=====] - 7s 8ms/step - loss: 0.5784 - accuracy: 0.9347 - val\_loss: 0.4616 - val\_accuracy: 0.9600

Epoch 5/50

860/860 [=====] - 7s 8ms/step - loss: 0.5354 - accuracy: 0.9382 - val\_loss: 0.4349 - val\_accuracy: 0.9650

Epoch 6/50

860/860 [=====] - 7s 8ms/step - loss: 0.5069 - accuracy: 0.9418 - val\_loss: 0.4030 - val\_accuracy: 0.9680

Epoch 7/50

860/860 [=====] - 7s 8ms/step - loss: 0.4922 - accuracy: 0.9414 - val\_loss: 0.4635 - val\_accuracy: 0.9474

Epoch 8/50

860/860 [=====] - 7s 8ms/step - loss: 0.4738 - accuracy: 0.9428 - val\_loss: 0.3959 - val\_accuracy: 0.9614

Epoch 9/50

860/860 [=====] - 7s 8ms/step - loss: 0.4534 - accuracy: 0.9450 - val\_loss: 0.4122 - val\_accuracy: 0.9580

Epoch 10/50

860/860 [=====] - 7s 8ms/step - loss: 0.4458 - accuracy: 0.9465 - val\_loss: 0.3582 - val\_accuracy: 0.9668

Epoch 11/50

860/860 [=====] - 7s 8ms/step - loss: 0.4350 - accuracy: 0.9471 - val\_loss: 0.3781 - val\_accuracy: 0.9628

Epoch 12/50

860/860 [=====] - 7s 9ms/step - loss: 0.4228 - accuracy: 0.9479 - val\_loss: 0.3549 - val\_accuracy: 0.9666

Epoch 13/50

860/860 [=====] - 7s 8ms/step - loss: 0.4117 - accuracy: 0.9487 - val\_loss: 0.3418 - val\_accuracy: 0.9684

Epoch 14/50

860/860 [=====] - 7s 8ms/step - loss: 0.4042 - accuracy: 0.9495 - val\_loss: 0.3588 - val\_accuracy: 0.9644

Epoch 15/50

860/860 [=====] - 7s 8ms/step - loss: 0.3965 - accuracy: 0.9505 - val\_loss: 0.3241 - val\_accuracy: 0.9702

Epoch 16/50

860/860 [=====] - 7s 8ms/step - loss: 0.3865 - accuracy: 0.9513 - val\_loss: 0.3574 - val\_accuracy: 0.9616

Epoch 17/50

860/860 [=====] - 7s 8ms/step - loss: 0.3853 - accuracy: 0.9512 - val\_loss: 0.3323 - val\_accuracy: 0.9684

Epoch 18/50

860/860 [=====] - 7s 8ms/step - loss: 0.3793 - accuracy: 0.9520 - val\_loss: 0.3197 - val\_accuracy: 0.9714

Epoch 19/50

860/860 [=====] - 7s 8ms/step - loss: 0.3799 - accuracy: 0.9504 - val\_loss: 0.3325 - val\_accuracy: 0.9670

Epoch 20/50

860/860 [=====] - 7s 8ms/step - loss: 0.3744 - accuracy: 0.9526 - val\_loss: 0.3292 - val\_accuracy: 0.9672

Epoch 21/50

860/860 [=====] - 7s 8ms/step - loss: 0.3701 - accuracy: 0.9527 - val\_loss: 0.3064 - val\_accuracy: 0.9724

Epoch 22/50

860/860 [=====] - 7s 8ms/step - loss: 0.3666 - accuracy: 0.9521 - val\_loss: 0.2888 - val\_accuracy: 0.9758

Epoch 23/50

860/860 [=====] - 7s 8ms/step - loss: 0.3605 - accuracy: 0.9539 - val\_loss: 0.2902 - val\_accuracy: 0.9744

Epoch 24/50

860/860 [=====] - 7s 8ms/step - loss: 0.3606 - accuracy: 0.9533 - val\_loss: 0.3204 - val\_accuracy: 0.9678

Epoch 25/50

860/860 [=====] - 7s 8ms/step - loss: 0.3580 - accuracy: 0.9539 - val\_loss: 0.3051 - val\_accuracy: 0.9698

Epoch 26/50

860/860 [=====] - 7s 8ms/step - loss: 0.3573 - accuracy: 0.9537 - val\_loss: 0.3073 - val\_accuracy: 0.9716

Epoch 27/50

860/860 [=====] - 7s 8ms/step - loss: 0.3508 - accuracy: 0.9555 - val\_loss: 0.3069 - val\_accuracy: 0.9692

Epoch 28/50


860/860 [=====] - 7s 8ms/step - loss: 0.3568 - accuracy: 0.9541 - val\_loss: 0.3036 - val\_accuracy: 0.9678

Epoch 29/50

860/860 [=====] - 7s 8ms/step - loss: 0.3457 - accuracy: 0.9571 - val\_loss: 0.2913 - val\_accuracy: 0.9734

```
In [38]: pd.DataFrame(history.history).plot()
```

Out[38]: <AxesSubplot: >



```
In [39]: ckpt_model = tf.keras.models.load_model(CKPT_path)
         ckpt_model.evaluate(x_test, y_test)
```

313/313 [=====] - 1s 2ms/step - loss: 0.3027 - accuracy: 0.9693

Out[39]: [0.30268532037734985, 0.9692999720573425]

The accuracy of the third **ANN** model is **96.93 %**

```
In [40]: del model
         del ckpt_model
```

## LeNet5 CNN Architecture

### Basic Introduction

LeNet-5, from the paper Gradient-Based Learning Applied to Document Recognition, is a very efficient convolutional neural network for handwritten character recognition.

#### Structure of the LeNet network

LeNet5 is a small network, it contains the basic modules of deep learning: convolutional layer, pooling layer, and full link layer. It is the basis of other deep learning models. Here we analyze LeNet5 in depth. At the same time, through example analysis, deepen the understanding of the convolutional layer and pooling layer.

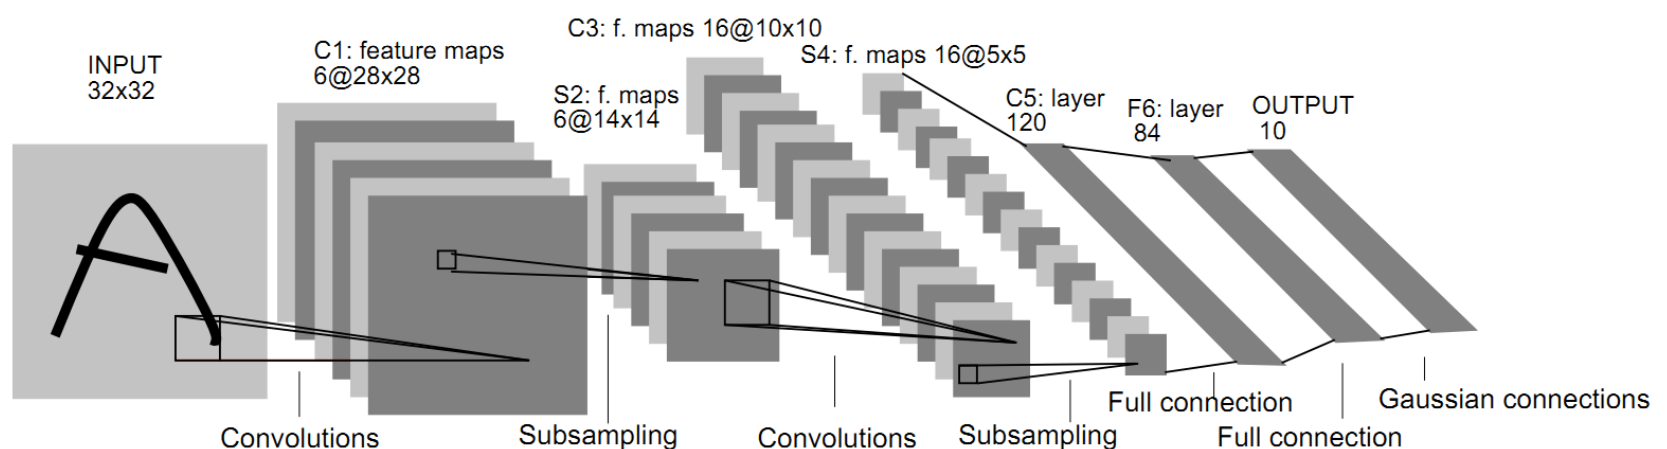


Fig. 2. Architecture of LeNet-5, a Convolutional Neural Network, here for digits recognition. Each plane is a feature map, i.e. a set of units whose weights are constrained to be identical.

LeNet-5 Total seven layer , does not comprise an input, each containing a trainable parameters; each layer has a plurality of the Map the Feature , a characteristic of each of the input FeatureMap extracted by means of a convolution filter, and then each FeatureMap There are multiple neurons.

Layer		Feature Map	Size	Kernel Size	Stride	Activation
Input	Image	1	32x32	-	-	-
1	Convolution	6	28x28	5x5	1	tanh
2	Average Pooling	6	14x14	2x2	2	tanh
3	Convolution	16	10x10	5x5	1	tanh
4	Average Pooling	16	5x5	2x2	2	tanh
5	Convolution	120	1x1	5x5	1	tanh
6	FC	-	84	-	-	tanh
Output	FC	-	10	-	-	softmax

Detailed explanation of each layer parameter:

INPUT Layer

The first is the data INPUT layer. The size of the input image is uniformly normalized to 32 \* 32.

Note: This layer does not count as the network structure of LeNet-5. Traditionally, the input layer is not considered as one of the network hierarchy.

C1 layer-convolutional layer

- Input picture:** 32 \* 32
- Convolution kernel size:** 5 \* 5
- Convolution kernel types:** 6
- Output featuremap size:** 28 \* 28 (32-5 + 1) = 28
- Number of neurons:** 28 \* 28 \* 6
- Trainable parameters:** (5 \* 5 + 1) \* 6 (5 \* 5 = 25 unit parameters and one bias parameter per filter, a total of 6 filters)
- Number of connections:** (5 \* 5 + 1) \* 6 \* 28 \* 28 = 122304

Detailed description:

- The first convolution operation is performed on the input image (using 6 convolution kernels of size 5 \* 5) to obtain 6 C1 feature maps (6 feature maps of size 28 \* 28, 32-5 + 1 = 28).
- Let's take a look at how many parameters are needed. The size of the convolution kernel is 5 \* 5, and there are 6 \* (5 \* 5 + 1) = 156 parameters in total, where +1 indicates that a kernel has a bias.
- For the convolutional layer C1, each pixel in C1 is connected to 5 \* 5 pixels and 1 bias in the input image, so there are 156 \* 28 \* 28 = 122304 connections in total. There are 122,304 connections, but we only need to learn 156 parameters, mainly through weight sharing.

S2 layer-pooling layer (downsampling layer)

- Input:** 28 \* 28
- Sampling area:** 2 \* 2
- Sampling method:** 4 inputs are added, multiplied by a trainable parameter, plus a trainable offset. Results via sigmoid
- Sampling type:** 6
- Output featureMap size:** 14 \* 14 (28/2)
- Number of neurons:** 14 \* 14 \* 6
- Trainable parameters:** 2 \* 6 (the weight of the sum + the offset)
- Number of connections:** (2 \* 2 + 1) \* 6 \* 14 \* 14
- The size of each feature map in S2 is 1/4 of the size of the feature map in C1.

Detailed description:

The pooling operation is followed immediately after the first convolution. Pooling is performed using 2 \* 2 kernels, and S2, 6 feature maps of 14 \* 14 (28/2 = 14) are obtained.

The pooling layer of S2 is the sum of the pixels in the 2 \* 2 area in C1 multiplied by a weight coefficient plus an offset, and then the result is mapped again.

So each pooling core has two training parameters, so there are  $2 \times 6 = 12$  training parameters, but there are  $5 \times 14 \times 14 \times 6 = 5880$  connections.


**C3 layer-convolutional layer**

- Input:** all 6 or several feature map combinations in S2
- Convolution kernel size:**  $5 * 5$
- Convolution kernel type:** 16
- Output featureMap size:**  $10 * 10$  ( $14 - 5 + 1$ ) = 10
- Each feature map in C3 is connected to all 6 or several feature maps in S2, indicating that the feature map of this layer is a different combination of the feature maps extracted from the previous layer.
- One way is that the first 6 feature maps of C3 take 3 adjacent feature map subsets in S2 as input. The next 6 feature maps take 4 subsets of neighboring feature maps in S2 as input. The next three take the non-adjacent 4 feature map subsets as input. The last one takes all the feature maps in S2 as input.
- The trainable parameters are:**  $6 * (3 * 5 * 5 + 1) + 6 * (4 * 5 * 5 + 1) + 3 * (4 * 5 * 5 + 1) + 1 * (6 * 5 * 5 + 1) = 1516$
- Number of connections:**  $10 * 10 * 1516 = 151600$


**Detailed description:**

After the first pooling, the second convolution, the output of the second convolution is C3, 16 10x10 feature maps, and the size of the convolution kernel is  $5 * 5$ . We know that S2 has 6  $14 * 14$  feature maps, how to get 16 feature maps from 6 feature maps? Here are the 16 feature maps calculated by the special combination of the feature maps of S2. details as follows:

The first 6 feature maps of C3 (corresponding to the 6th column of the first red box in the figure above) are connected to the 3 feature maps connected to the S2 layer (the first red box in the above figure), and the next 6 feature maps are connected to the S2 layer The 4 feature maps are connected (the second red box in the figure above), the next 3 feature maps are connected with the 4 feature maps that are not connected at the S2 layer, and the last is connected with all the feature maps at the S2 layer. The convolution kernel size is still  $5 * 5$ , so there are  $6 * (3 * 5 * 5 + 1) + 6 * (4 * 5 * 5 + 1) + 3 * (4 * 5 * 5 + 1) + 1 * (6 * 5 * 5 + 1) = 1516$  parameters. The image size is  $10 * 10$ , so there are 151600 connections.



The convolution structure of C3 and the first 3 graphs in S2 is shown below:



**S4 layer-pooling layer (downsampling layer)**

- Input:**  $10 * 10$
- Sampling area:**  $2 * 2$
- Sampling method:** 4 inputs are added, multiplied by a trainable parameter, plus a trainable offset. Results via sigmoid

**Sampling type:** 16

**Output featureMap size:**  $5 * 5 (10/2)$

**Number of neurons:**  $5 * 5 * 16 = 400$

**Trainable parameters:**  $2 * 16 = 32$  (the weight of the sum + the offset)

**Number of connections:**  $16 * (2 * 2 + 1) * 5 * 5 = 2000$

The size of each feature map in S4 is 1/4 of the size of the feature map in C3

#### Detailed description:

S4 is the pooling layer, the window size is still  $2 * 2$ , a total of 16 feature maps, and the 16  $10 \times 10$  maps of the C3 layer are pooled in units of  $2 \times 2$  to obtain 16  $5 \times 5$  feature maps. This layer has a total of 32 training parameters of  $2 \times 16$ ,  $5 \times 5 \times 16 = 2000$  connections.

*The connection is similar to the S2 layer.*

### C5 layer-convolution layer

**Input:** All 16 unit feature maps of the S4 layer (all connected to s4)

**Convolution kernel size:**  $5 * 5$


**Convolution kernel type:** 120

**Output featureMap size:**  $1 * 1 (5-5 + 1)$

**Trainable parameters / connection:**  $120 * (16 * 5 * 5 + 1) = 48120$

#### Detailed description:

The C5 layer is a convolutional layer. Since the size of the 16 images of the S4 layer is  $5 \times 5$ , which is the same as the size of the convolution kernel, the size of the image formed after convolution is  $1 \times 1$ . This results in 120 convolution results. Each is connected to the 16 maps on the previous level. So there are  $(5 \times 5 \times 16 + 1) \times 120 = 48120$  parameters, and there are also 48120 connections. The network structure of the C5 layer is as follows:



### F6 layer-fully connected layer

**Input:** c5 120-dimensional vector

**Calculation method:** calculate the dot product between the input vector and the weight vector, plus an offset, and the result is output through the sigmoid function.

**Trainable parameters:**  $84 * (120 + 1) = 10164$

#### Detailed description:

Layer 6 is a fully connected layer. The F6 layer has 84 nodes, corresponding to a  $7 \times 12$  bitmap, -1 means white, 1 means black, so the black and white of the bitmap of each symbol corresponds to a code. The training parameters and number of connections for this layer are  $(120 + 1) \times 84 = 10164$ . The ASCII encoding diagram is as follows:



The connection method of the F6 layer is as follows:



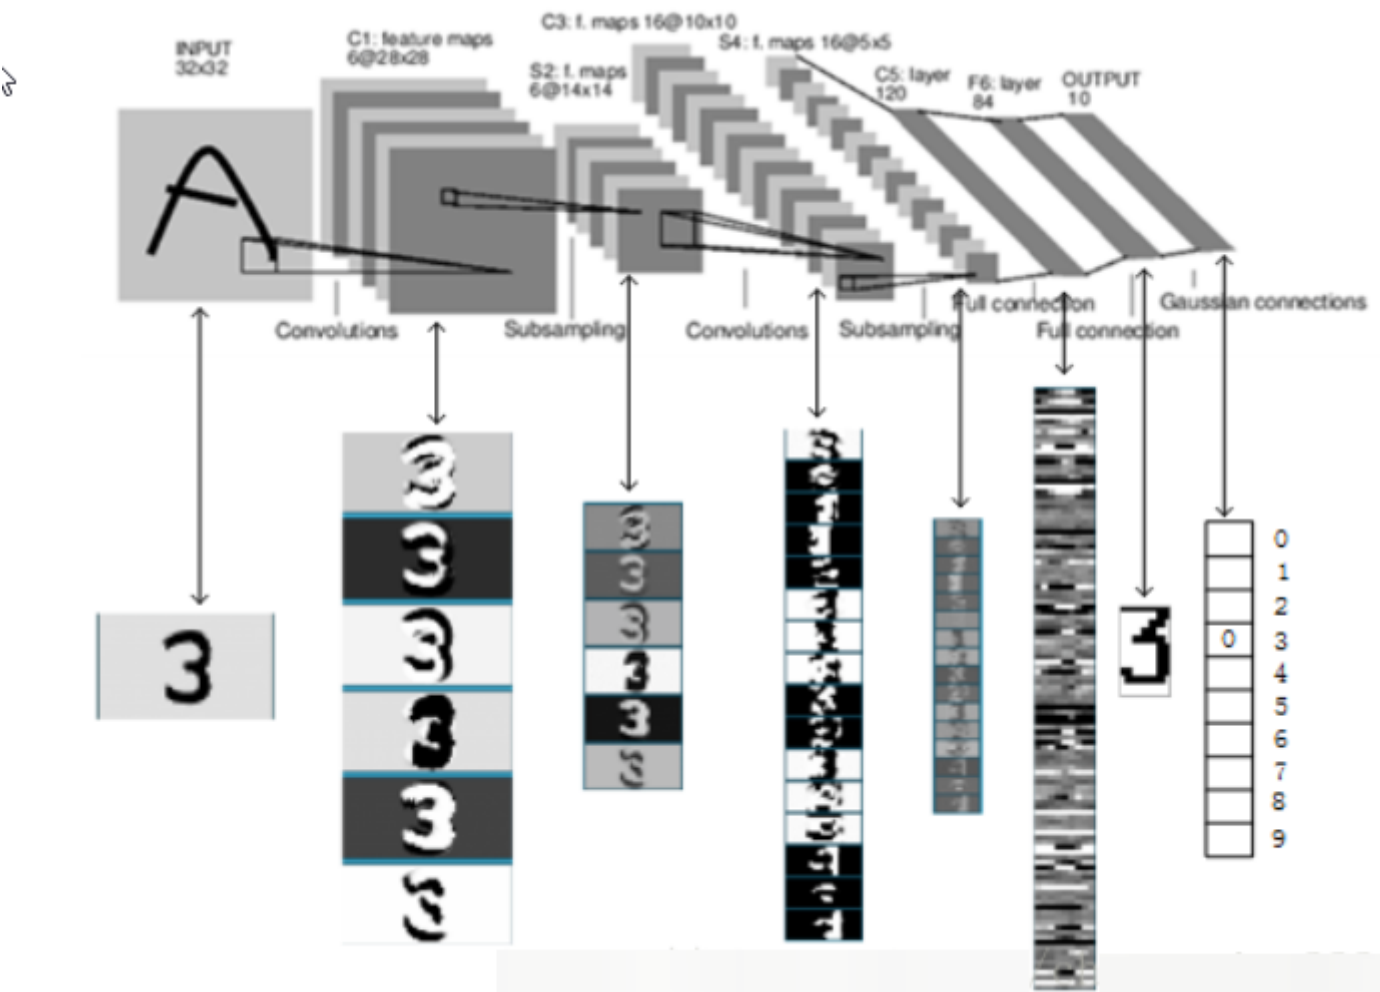
### Output layer-fully connected layer

The output layer is also a fully connected layer, with a total of 10 nodes, which respectively represent the numbers 0 to 9, and if the value of node  $i$  is 0, the result of network recognition is the number  $i$ . A radial basis function (RBF) network connection is used. Assuming  $x$  is the input of the previous layer and  $y$  is the output of the RBF, the calculation of the RBF output is:

$$y_i = \sum_j (x_j - w_{ij})^2$$

The value of the above formula  $w_{ij}$  is determined by the bitmap encoding of  $i$ , where  $i$  ranges from 0 to 9, and  $j$  ranges from 0 to  $7 * 12 - 1$ .

The closer the value of the RBF output is to 0, the closer it is to  $i$ , that is, the closer to the ASCII encoding figure of  $i$ , it means that the recognition result input by the current network is the character  $i$ . This layer has  $84 \times 10 = 840$  parameters and connections.



### Summary

- LeNet-5 is a very efficient convolutional neural network for handwritten character recognition.
- Convolutional neural networks can make good use of the structural information of images.
- The convolutional layer has fewer parameters, which is also determined by the main characteristics of the convolutional layer, that is, local connection and shared weights.

```
In [41]: # Reshaping the data so that it is compatible with the architecture
x_train.reshape(x_train.shape[0],x_train.shape[1],x_train.shape[2],1)
x_test.reshape(x_test.shape[0],x_test.shape[1],x_test.shape[2],1)
```

```
# Building the Model Architecture
model = Sequential()

model.add(Conv2D(6, kernel_size = (5,5), padding = 'valid', activation='tanh', input_shape = (28,28,1)))
model.add(AveragePooling2D(pool_size=(2,2), strides = 2, padding = 'valid'))

model.add(Conv2D(16, kernel_size = (5,5), padding = 'valid', activation='tanh'))
model.add(AveragePooling2D(pool_size=(2,2), strides = 2, padding = 'valid'))

model.add(Flatten())

model.add(Dense(120, activation='tanh'))
model.add(Dense(84, activation='tanh'))
model.add(Dense(10, activation='softmax'))
```

In [42]: `model.summary()`

Model: "sequential\_3"

Layer (type)	Output Shape	Param #
=====		
conv2d (Conv2D)	(None, 24, 24, 6)	156
average_pooling2d (Average Pooling2D)	(None, 12, 12, 6)	0
conv2d_1 (Conv2D)	(None, 8, 8, 16)	2416
average_pooling2d_1 (Average Pooling2D)	(None, 4, 4, 16)	0
flatten_2 (Flatten)	(None, 256)	0
dense_12 (Dense)	(None, 120)	30840
dense_13 (Dense)	(None, 84)	10164
dense_14 (Dense)	(None, 10)	850
=====		
Total params: 44426 (173.54 KB)		
Trainable params: 44426 (173.54 KB)		
Non-trainable params: 0 (0.00 Byte)		

```
In [43]: def get_log_path(log_dir="logs/fit"):
    filename = time.strftime("4_log_%y_%m_%d_%H_%M_%S")
    logs_path = os.path.join(log_dir, filename)
    print(f"Saving logs at {logs_path}")
    return logs_path

log_dirs = get_log_path()
tb_cb = tf.keras.callbacks.TensorBoard(log_dir=log_dirs)

early_stopping_cb = tf.keras.callbacks.EarlyStopping(patience=7, restore_best_weights=True)

CKPT_path = os.path.join("Models", "Model_ckpt_Digit_mnist_4.h5")
checkpoint_cb = tf.keras.callbacks.ModelCheckpoint(CKPT_path, save_best_only=True)

EPOCHS = 50
VALIDATION_SET = (x_valid, y_valid)

loss_function = "sparse_categorical_crossentropy"
OPTIMIZER = tf.keras.optimizers.Adam(learning_rate=0.001)
METRICS = ["accuracy"]

model.compile(
    loss=loss_function,
    optimizer=OPTIMIZER,
    metrics=METRICS
)

history = model.fit(x_train, y_train, epochs=EPOCHS, validation_data=VALIDATION_SET, batch_size=64, verbose=1,
    callbacks=[tb_cb, early_stopping_cb, checkpoint_cb], use_multiprocessing=True)
```



Saving logs at logs/fit/4\_log\_23\_09\_26\_19\_58\_56

Epoch 1/50

860/860 [=====] - 8s 8ms/step - loss: 0.2934 - accuracy: 0.9136 - val\_loss: 0.1174 - val\_accuracy: 0.9664

Epoch 2/50

860/860 [=====] - 7s 8ms/step - loss: 0.1076 - accuracy: 0.9671 - val\_loss: 0.0772 - val\_accuracy: 0.9756

Epoch 3/50

860/860 [=====] - 7s 8ms/step - loss: 0.0720 - accuracy: 0.9780 - val\_loss: 0.0634 - val\_accuracy: 0.9820

Epoch 4/50

860/860 [=====] - 7s 8ms/step - loss: 0.0563 - accuracy: 0.9827 - val\_loss: 0.0515 - val\_accuracy: 0.9848

Epoch 5/50

860/860 [=====] - 7s 8ms/step - loss: 0.0439 - accuracy: 0.9861 - val\_loss: 0.0490 - val\_accuracy: 0.9852

Epoch 6/50

860/860 [=====] - 7s 8ms/step - loss: 0.0365 - accuracy: 0.9883 - val\_loss: 0.0482 - val\_accuracy: 0.9860

Epoch 7/50

860/860 [=====] - 7s 8ms/step - loss: 0.0309 - accuracy: 0.9901 - val\_loss: 0.0447 - val\_accuracy: 0.9864

Epoch 8/50

860/860 [=====] - 7s 8ms/step - loss: 0.0258 - accuracy: 0.9919 - val\_loss: 0.0545 - val\_accuracy: 0.9832

Epoch 9/50

860/860 [=====] - 7s 8ms/step - loss: 0.0214 - accuracy: 0.9930 - val\_loss: 0.0407 - val\_accuracy: 0.9884

Epoch 10/50

860/860 [=====] - 7s 8ms/step - loss: 0.0180 - accuracy: 0.9943 - val\_loss: 0.0580 - val\_accuracy: 0.9842

Epoch 11/50

860/860 [=====] - 7s 8ms/step - loss: 0.0169 - accuracy: 0.9945 - val\_loss: 0.0430 - val\_accuracy: 0.9888

Epoch 12/50

860/860 [=====] - 7s 8ms/step - loss: 0.0133 - accuracy: 0.9954 - val\_loss: 0.0503 - val\_accuracy: 0.9864

Epoch 13/50

860/860 [=====] - 7s 8ms/step - loss: 0.0122 - accuracy: 0.9964 - val\_loss: 0.0469 - val\_accuracy: 0.9876

Epoch 14/50

860/860 [=====] - 7s 8ms/step - loss: 0.0107 - accuracy: 0.9964 - val\_loss: 0.0416 - val\_accuracy: 0.9898

Epoch 15/50


860/860 [=====] - 7s 8ms/step - loss: 0.0100 - accuracy: 0.9968 - val\_loss: 0.0464 - val\_accuracy: 0.9872

Epoch 16/50

860/860 [=====] - 7s 8ms/step - loss: 0.0078 - accuracy: 0.9976 - val\_loss: 0.0475 - val\_accuracy: 0.9880

```
In [44]: pd.DataFrame(history.history).plot()
```

Out[44]: <AxesSubplot: >



```
In [45]: ckpt_model = tf.keras.models.load_model(CKPT_path)
         ckpt_model.evaluate(x_test, y_test)
```

313/313 [=====] - 1s 3ms/step - loss: 0.0432 - accuracy: 0.9861

Out[45]: [0.04321447014808655, 0.9861000180244446]

The accuracy of the **CNN** model is **98.61 %**



```
In [46]: # %load_ext tensorboard
# %tensorboard --logdir logs/fit
```

*Thank You*